

Team Details

Team Name: EMBEDDEDX

SR. NO	ROLE	NAME	ACADEMIC YEAR
1	Team Leader	KULDEEP SHARMA	2ND YEAR
2	Member 1	PRAFUL KUMAR SINGH	2ND YEAR
3	Member 2	ANURAG SAHOO	3RD YEAR
4	Member 3	PRITHVI RAJ	3RD YEAR

 A team can have up to 4 members including the team leader. Add rows if necessary.

COLLEGE NAME

INDIAN INSTITUTE OF INFORMATION TECHNOLOGY GUWAHATI

TEAM LEADER CONTACT NUMBER

+91 7014208868

TEAM LEADER EMAIL ADDRESS

kuldeep.sharma24b@iiitg.ac.in

Problem Statement Addressed

In semiconductor manufacturing, wafer surface inspection is a critical quality control step. Currently, this process relies heavily on manual visual inspection by human operators using Scanning Electron Microscope (SEM) images. This approach suffers from three major limitations:

- 1) Slow — Human inspectors can only analyze a limited number of wafers per hour, creating production bottlenecks as fabrication volumes scale up.
- 2) Error-Prone — Fatigue, subjectivity, and the subtle nature of microscopic defects (cracks, pits, particle contamination, scratches) lead to inconsistent detection rates and missed defects, resulting in defective chips reaching downstream processes.
- 3) Not Scalable — As semiconductor nodes shrink (7nm, 5nm, 3nm), defect patterns become increasingly complex. Training new inspectors is expensive and slow, and human throughput cannot match fab output demands.

There is an urgent need for a lightweight, accurate AI model that can perform real-time defect classification directly on edge devices (embedded in inspection equipment on the fab floor), eliminating dependency on cloud connectivity, reducing latency, and enabling 24/7 automated inspection across 10 distinct defect categories including cracks, scratches, particle contamination, pits/voids, and more

Idea Description

We propose a Hybrid CNN + SVM architecture that combines the powerful feature extraction capability of a fine-tuned SqueezeNet (a lightweight CNN with only ~0.7M parameters) with the robust classification ability of a Support Vector Machine (SVM) with RBF kernel.

Instead of using SqueezeNet's built-in classifier, we strip the classifier head and use the CNN purely as a 512-dimensional feature extractor. These rich, learned features are then fed into an SVM, which draws optimal decision boundaries between the 10 defect classes. This hybrid approach outperforms both standalone CNN and standalone SVM solutions.

Solution Overview : The solution directly addresses the problem by:

- 1)Speed — SqueezeNet is 50× smaller than AlexNet and designed for edge deployment. Feature extraction + SVM prediction takes milliseconds per image.
- 2)Accuracy — The hybrid model achieves 98.54% accuracy across all 10 SEM defect classes with a 0.96 macro F1-score, meaning it performs consistently well even on rare defect types.
- 3)Scalability — The entire deployment package is just 3 small files (CNN weights ~3MB + SVM model ~1MB + class labels <1KB), deployable on embedded devices, Raspberry Pi, or NVIDIA Jetson without cloud dependency.
- 4)Robustness — SVM's maximum-margin classification is more resistant to overfitting on small per-class samples compared to a softmax classifier.

Proposed Solution

Methodology: The solution follows a 3-stage pipeline:

Stage 1 — Fine-Tuning SqueezeNet (Transfer Learning) :-

Pre-trained SqueezeNet1.1 (ImageNet weights) is fine-tuned on 4,352 SEM wafer images across 10 defect classes

Training uses Adam optimizer (lr=5e-4), CosineAnnealing LR scheduler, weight decay (1e-4), and data augmentation (random flip, rotation, crop) to prevent overfitting

Best model checkpoint achieves 98.99% validation accuracy in 30 epochs

Stage 2 — Feature Extraction :-

The classifier head is stripped; only the convolutional feature layers are retained

Each SEM image is passed through the CNN to produce a 512-dimensional feature vector via Global Average Pooling

This converts raw pixels into semantically rich, compact representations

Stage 3 — SVM Classification :-

An SVM with RBF kernel (C=10, gamma='scale') is trained on the 512-dim CNN features

SVM finds the optimal hyperplane separating 10 defect classes in high-dimensional feature space

Technologies:

1)PyTorch (CNN training & inference)

2)scikit-learn (SVM classifier)

3)SqueezeNet1.1 architecture (edge-optimized, ~0.7M parameters)

4)Deployed as 3 files: CNN weights (.pth) + SVM model (.pkl) + class labels (.pkl)

Implementation Strategy:

Trained on GPU (Kaggle/cloud) → Export 3 lightweight files → Deploy on edge device

Inference pipeline: SEM Image → Resize 224×224 → CNN Features (512-dim) → SVM Predict → Defect Class

Total inference time: <50ms per image on edge hardware

Results: 98.54% accuracy | 0.96 macro F1 | All 10 classes correctly classified

I Innovation and Uniqueness

○ Innovation : Key Innovation:

Our solution introduces a novel Hybrid CNN+SVM architecture for semiconductor defect inspection. Instead of using SqueezeNet as a full classifier, we strip its classifier head and use only the convolutional layers as a 512-dimensional feature extractor. These features are then classified by an SVM with RBF kernel, which finds mathematically optimal decision boundaries. This hybrid approach is a new application in SEM wafer inspection. Additionally, we bridge the domain gap from natural images (ImageNet) to grayscale SEM microscopy through careful transfer learning with augmentation — proving this technique works for specialized industrial imaging.

Competitive Advantage:

Our model is 50× smaller than AlexNet and 130× smaller than VGG, totaling just ~4MB for full deployment. It achieves 98.54% accuracy compared to ~95-97% for CNN-only approaches on SEM data. SVM's maximum-margin principle handles rare defect classes better than softmax, giving us a strong 0.96 macro F1-score across all 10 classes. The entire system runs in under 50ms on CPU-only edge devices like Raspberry Pi (\$35), requires zero cloud connectivity, and deploys as just 3 files — making it the most cost-effective, lightweight, and accurate solution compared to existing heavy-model alternatives that need expensive GPU hardware and cloud infrastructure.

■ Impact and Benefits

○ Primary Impact:

Eliminates manual wafer inspection bottleneck in semiconductor fabs. Our hybrid SqueezeNet + SVM model replaces slow, inconsistent human visual inspection with an automated AI system that classifies 10 types of SEM wafer defects in real-time, directly on the production floor. The system operates 24/7 with zero fatigue errors, catches defective wafers instantly before they contaminate downstream fabrication steps, and works fully offline with no cloud dependency — ensuring complete data privacy of sensitive fab images.

Quantifiable Outcomes of using such a model:

98.54% classification accuracy (vs ~85-90% human average)

600× faster inspection — under 50ms per image vs 30-60 seconds for human inspectors

99% cost reduction per inspection — ~₹0.01/image compute cost vs ₹500-1000/hr skilled labor

4MB total model size — deployable on \$35 edge hardware

0.96 macro F1-score — balanced performance across all 10 defect classes including rare ones

30-minute training time on GPU vs 3-6 months to train a new human inspector

Zero downtime — continuous 24/7 automated operation with consistent accuracy

Technology & Feasibility/Methodology Used



Technology and Feasibility : Implementation Strategy:

The system follows a Train-on-Cloud, Deploy-on-Edge architecture. SqueezeNet 1.1 (pre-trained on ImageNet) is fine-tuned on 4,352 SEM wafer images using PyTorch with Adam optimizer, CosineAnnealing LR scheduler, and data augmentation to prevent overfitting. After 30 epochs of training on Kaggle GPU, the classifier head is stripped and the convolutional layers serve as a 512-dim feature extractor. An SVM with RBF kernel ($C=10$, $\gamma=\text{scale}$) is trained on these features using scikit-learn. The final deployment package is just 3 files (~4MB): CNN weights (.pth), SVM model (.pkl), and class labels (.pkl). The inference pipeline is fully offline — SEM image goes in, defect class comes out in under 50ms. The approach is validated with 98.54% test accuracy on real SEM data and is hardware-agnostic.

Software Architecture:

PyTorch for deep learning training and inference. SqueezeNet 1.1 as the lightweight CNN backbone with only 0.7M parameters. scikit-learn for SVM classification with RBF kernel. torchvision for image preprocessing, normalization, and data augmentation. Pillow (PIL) for handling 8-bit, 16-bit, and grayscale SEM image formats. joblib for serializing SVM model and class labels. matplotlib and seaborn for confusion matrix and training curve visualization.

Hardware Components:

Training Phase — NVIDIA T4/P100 GPU via Kaggle with 16GB VRAM and 50GB storage for dataset and checkpoints.

Development Tools:

Kaggle Notebooks for GPU-accelerated training. Python 3.12 as the primary language. PyTorch 2.x for deep learning. scikit-learn for SVM. Git/GitHub for version control. ONNX Runtime and TensorRT planned for future optimized edge inference.

GitHub & Drive Links



GitHub Repository

<https://github.com/KuldeepSharma9351/IESA-Hackathon-EmbeddedX>

ONNX Model Link :

https://drive.google.com/drive/folders/1xP0DqTMVuJmFWPtpVLaqKSAr7FsG5_Gv

Dataset ZIP Link:

<https://drive.google.com/drive/folders/1XN-VkU2b6Ix9Z77EBKR0vI5C7wzfyAPg>

Results Report Link:

<https://drive.google.com/drive/folders/1qISNLv7OEgMYurCHuMA1k4z5uuX7OAS7>

Research and References



Research Background & Methodology

Semiconductor wafer defect detection is a critical quality control challenge. Traditional manual inspection of SEM images is subjective, slow, and inconsistent. Recent research has demonstrated that deep learning, particularly CNNs, can significantly improve defect classification accuracy and speed.

Our work builds upon three key research directions. First, deep CNNs have been shown to outperform manual classification for wafer defect recognition by automatically learning discriminative features from SEM images, eliminating human subjectivity and improving reliability. Second, lightweight architectures like SqueezeNet have proven effective for semiconductor defect classification when combined with geometric transformation-based data augmentation, achieving competitive accuracy with 50× fewer parameters than traditional deep models — making edge deployment feasible. Third, recent work on long-tailed wafer defect detection from SEM images has highlighted the challenges of imbalanced defect distributions, diverse image backgrounds, and varying defect scales — problems common in real fab environments.

Our methodology addresses all three challenges by fine-tuning a pre-trained SqueezeNet on SEM wafer images with data augmentation (random flip, rotation, crop) to handle limited and imbalanced training data. We then strip the CNN classifier and extract 512-dimensional feature vectors, which are classified using an SVM with RBF kernel. This hybrid CNN+SVM approach leverages the CNN's feature learning capability with SVM's maximum-margin classification principle, achieving 98.54% accuracy across 10 defect classes with a strong 0.96 macro F1-score, demonstrating robustness even on underrepresented defect categories.



References & Citations

List key papers, articles, or data sources.

<https://github.com/SpatialAILab/WaferDC>

<https://ieeexplore.ieee.org/document/8578572>

<https://ieeexplore.ieee.org/document/10448555>

<https://www.sciencedirect.com/science/article/pii/S0360835223005739>